

# Continuous\_Control

July 17, 2019

## 1 Continuous Control

---

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

### 1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [8]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible.  
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 2.0.2.
```

The environments corresponding to both versions of the environment are already saved in the Workspace and can be accessed at the file paths provided below.

Please select one of the two options below for loading the environment.

```
In [9]: from unityagents import UnityEnvironment  
import numpy as np  
  
# select this option to load version 1 (with a single agent) of the environment  
#env = UnityEnvironment(file_name='/data/Reacher_One_Linux_NoVis/Reacher_One_Linux_NoVis')  
env = UnityEnvironment(file_name='/data/Reacher_Linux_NoVis/Reacher.x86_64')  
# select this option to load version 2 (with 20 agents) of the environment  
# env = UnityEnvironment(file_name='/data/Reacher_Linux_NoVis/Reacher.x86_64')
```

```
INFO:unityagents:  
'Academy' started successfully!  
Unity Academy name: Academy  
Number of Brains: 1  
Number of External Brains : 1  
Lesson number : 0  
Reset Parameters :  
    goal_speed -> 1.0
```

```

        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 33
    Number of stacked Vector Observation: 1
    Vector Action space type: continuous
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [11]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

## 1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [12]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents
        num_agents = len(env_info.agents)
        print('Number of agents:', num_agents)

        # size of each action
        action_size = brain.vector_action_space_size
        print('Size of each action:', action_size)

        # examine the state space
        states = env_info.vector_observations
        state_size = states.shape[1]
        print('There are {} agents. Each observes a state with length: {}'.format(states.shape[0], state_size))
        print('The state for the first agent looks like:', states[0])

```

Number of agents: 20

Size of each action: 4

There are 20 agents. Each observes a state with length: 33

The state for the first agent looks like: [ 0.00000000e+00 -4.00000000e+00 0.00000000e+00  
-0.00000000e+00 -0.00000000e+00 -4.37113883e-08 0.00000000e+00  
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00  
0.00000000e+00 0.00000000e+00 -1.00000000e+01 0.00000000e+00  
1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08  
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00

```

0.00000000e+00  0.00000000e+00  5.75471878e+00 -1.00000000e+00
5.55726624e+00  0.00000000e+00  1.00000000e+00  0.00000000e+00
-1.68164849e-01]

```

```

In [13]: env_info = env.reset(train_mode=True)[brain_name]      # reset the environment
          states = env_info.vector_observations                  # get the current state (for each agent)
          scores = np.zeros(num_agents)                         # initialize the score (for each agent)
          while True:
              actions = np.random.randn(num_agents, action_size) # select an action (for each agent)
              actions = np.clip(actions, -1, 1)                  # all actions between -1 and 1
              env_info = env.step(actions)[brain_name]           # send all actions to the environment
              next_states = env_info.vector_observations           # get next state (for each agent)
              rewards = env_info.rewards                          # get reward (for each agent)
              dones = env_info.local_done                        # see if episode finished
              scores += env_info.rewards                         # update the score (for each agent)
              states = next_states                               # roll over states to next time
              if np.any(dones):                                  # exit loop if episode finished
                  break
          print('Total score (averaged over agents) this episode: {}'.format(np.mean(scores)))

```

```
Total score (averaged over agents) this episode: 0.17149999616667627
```

### 1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agents while they are training**, and you should set `train_mode=True` to restart the environment.

When finished, you can close the environment.

```

In [14]: print(actions.shape)
          print(next_states.shape)
          print(len(rewards))

```

```

(20, 4)
(20, 33)
20

```

### 1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agents while they are training. However, *after training the agents*, you can download the saved model weights to watch the agents on your own machine!

```
In [16]: import numpy as np

import torch
import torch.nn as nn
import torch.nn.functional as F

def hidden_init(layer):
    fan_in = layer.weight.data.size()[0]
    lim = 1. / np.sqrt(fan_in)
    return (-lim, lim)

class Actor(nn.Module):
    """Actor (Policy) Model."""

    def __init__(self, state_size, action_size, seed=0, fc1_units=128, fc2_units=128):
        """Initialize parameters and build model.
        Params
        =====
            state_size (int): Dimension of each state
            action_size (int): Dimension of each action
            seed (int): Random seed
            fc1_units (int): Number of nodes in first hidden layer
            fc2_units (int): Number of nodes in second hidden layer
        """
        super(Actor, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fc1 = nn.Linear(state_size, fc1_units)
        self.bn1 = nn.BatchNorm1d(fc1_units)
        self.fc2 = nn.Linear(fc1_units, fc2_units)
        self.fc3 = nn.Linear(fc2_units, action_size)
        self.reset_parameters()

    def reset_parameters(self):
        self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state):
        """Build an actor (policy) network that maps states -> actions."""
        if state.dim() == 1:
            state = torch.unsqueeze(state, 0)
```

```

        x = F.relu(self.fc1(state))
        x = self.bn1(x)
        x = F.relu(self.fc2(x))
        return F.tanh(self.fc3(x))

class Critic(nn.Module):
    """Critic (Value) Model."""

    def __init__(self, state_size, action_size, seed=0, fcs1_units=128, fc2_units=128):
        """Initialize parameters and build model.
        Params
        =====
            state_size (int): Dimension of each state
            action_size (int): Dimension of each action
            seed (int): Random seed
            fcs1_units (int): Number of nodes in the first hidden layer
            fc2_units (int): Number of nodes in the second hidden layer
        """
        super(Critic, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fcs1 = nn.Linear(state_size, fcs1_units)
        self.bn1 = nn.BatchNorm1d(fcs1_units)
        self.fc2 = nn.Linear(fcs1_units+action_size, fc2_units)
        self.fc3 = nn.Linear(fc2_units, 1)
        self.reset_parameters()

    def reset_parameters(self):
        self.fcs1.weight.data.uniform_(*hidden_init(self.fcs1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state, action):
        """Build a critic (value) network that maps (state, action) pairs -> Q-values."""
        if state.dim() == 1:
            state = torch.unsqueeze(state,0)
        xs = F.relu(self.fcs1(state))
        xs = self.bn1(xs)
        x = torch.cat((xs, action), dim=1)
        x = F.relu(self.fc2(x))
        return self.fc3(x)

```

```

In [19]: import numpy as np
import random
import copy
from collections import namedtuple, deque

##from model import Actor, Critic

```

```

import torch
import torch.nn.functional as F
import torch.optim as optim

BUFFER_SIZE = int(1e6) # replay buffer size
BATCH_SIZE = 128      # minibatch size
GAMMA = 0.99          # discount factor
TAU = 1e-3            # for soft update of target parameters
LR_ACTOR = 1e-4        # learning rate of the actor
LR_CRITIC = 1e-4       # learning rate of the critic
WEIGHT_DECAY = 0.0     # L2 weight decay

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

class Agent():
    """Interacts with and learns from the environment."""

    def __init__(self, state_size, action_size, random_seed):
        """Initialize an Agent object.

        Params
        =====
            state_size (int): dimension of each state
            action_size (int): dimension of each action
            random_seed (int): random seed
        """
        self.state_size = state_size
        self.action_size = action_size
        self.seed = random.seed(random_seed)

        # Actor Network (w/ Target Network)
        self.actor_local = Actor(state_size, action_size, random_seed).to(device)
        self.actor_target = Actor(state_size, action_size, random_seed).to(device)
        self.actor_optimizer = optim.Adam(self.actor_local.parameters(), lr=LR_ACTOR)

        # Critic Network (w/ Target Network)
        self.critic_local = Critic(state_size, action_size, random_seed).to(device)
        self.critic_target = Critic(state_size, action_size, random_seed).to(device)
        self.critic_optimizer = optim.Adam(self.critic_local.parameters(), lr=LR_CRITIC)

        # Noise process
        self.noise = OUNoise(action_size, random_seed)

        # Replay memory
        self.memory = ReplayBuffer(action_size, BUFFER_SIZE, BATCH_SIZE, random_seed)

    def step(self, state, action, reward, next_state, done):
        """Save experience in replay memory, and use random sample from buffer to learn

```

```

# Save experience / reward
self.memory.add(state, action, reward, next_state, done)

# Learn, if enough samples are available in memory
#Commented out to improve performance. Now we learn only when start_learn is called
#if len(self.memory) > BATCH_SIZE:
    #experiences = self.memory.sample()
    #self.learn(experiences, GAMMA)

def act(self, state, add_noise=True):
    """Returns actions for given state as per current policy."""
    state = torch.from_numpy(state).float().to(device)
    self.actor_local.eval()
    with torch.no_grad():
        action = self.actor_local(state).cpu().data.numpy()
    self.actor_local.train()
    if add_noise:
        action += self.noise.sample()
    return np.clip(action, -1, 1)

def reset(self):
    self.noise.reset()

def start_learn(self):
    if len(self.memory) > BATCH_SIZE:
        experiences = self.memory.sample()
        self.learn(experiences, GAMMA)

def learn(self, experiences, gamma):
    """Update policy and value parameters using given batch of experience tuples.
    Q_targets = r +  $\gamma$  * critic_target(next_state, actor_target(next_state))
    where:
        actor_target(state) -> action
        critic_target(state, action) -> Q-value
    Params
    =====
        experiences (Tuple[torch.Tensor]): tuple of (s, a, r, s', done) tuples
        gamma (float): discount factor
    """
    states, actions, rewards, next_states, dones = experiences

    # ----- update critic ----- #
    # Get predicted next-state actions and Q values from target models
    actions_next = self.actor_target(next_states)
    Q_targets_next = self.critic_target(next_states, actions_next)
    # Compute Q targets for current states (y_i)
    Q_targets = rewards + (gamma * Q_targets_next * (1 - dones))
    # Compute critic loss

```

```

Q_expected = self.critic_local(states, actions)
critic_loss = F.mse_loss(Q_expected, Q_targets)
# Minimize the loss
self.critic_optimizer.zero_grad()
critic_loss.backward()
self.critic_optimizer.step()

# ----- update actor ----- #
# Compute actor loss
actions_pred = self.actor_local(states)
actor_loss = -self.critic_local(states, actions_pred).mean()
# Minimize the loss
self.actor_optimizer.zero_grad()
actor_loss.backward()
self.actor_optimizer.step()

# ----- update target networks ----- #
self.soft_update(self.critic_local, self.critic_target, TAU)
self.soft_update(self.actor_local, self.actor_target, TAU)

def soft_update(self, local_model, target_model, tau):
    """Soft update model parameters.
    _target = *_local + (1 - )*_target
    Params
    =====
        local_model: PyTorch model (weights will be copied from)
        target_model: PyTorch model (weights will be copied to)
        tau (float): interpolation parameter
    """
    for target_param, local_param in zip(target_model.parameters(), local_model.par
        target_param.data.copy_(tau*local_param.data + (1.0-tau)*target_param.data)

class OUNoise:
    """Ornstein-Uhlenbeck process."""

    def __init__(self, size, seed, mu=0., theta=0.15, sigma=0.2):
        """Initialize parameters and noise process."""
        self.mu = mu * np.ones(size)
        self.theta = theta
        self.sigma = sigma
        self.size = size
        self.seed = random.seed(seed)
        self.reset()

    def reset(self):
        """Reset the internal state (= noise) to mean (mu)."""
        self.state = copy.copy(self.mu)

```



```

def sample(self):
    """Update internal state and return it as a noise sample."""
    x = self.state

    # Thanks to Hiu C. for this tip, this really helped get the learning up to the
    dx = self.theta * (self.mu - x) + self.sigma * np.random.standard_normal(self.s

    self.state = x + dx
    return self.state

class ReplayBuffer:
    """Fixed-size buffer to store experience tuples."""

    def __init__(self, action_size, buffer_size, batch_size, seed):
        """Initialize a ReplayBuffer object.

        Params
        =====
            buffer_size (int): maximum size of buffer
            batch_size (int): size of each training batch
        """
        self.action_size = action_size
        self.memory = deque(maxlen=buffer_size) # internal memory (deque)
        self.batch_size = batch_size
        self.experience = namedtuple("Experience", field_names=["state", "action", "reward", "next_state", "done"])
        self.seed = random.seed(seed)

    def add(self, state, action, reward, next_state, done):
        """Add a new experience to memory."""
        e = self.experience(state, action, reward, next_state, done)
        self.memory.append(e)

    def sample(self):
        """Randomly sample a batch of experiences from memory."""
        experiences = random.sample(self.memory, k=self.batch_size)

        states = torch.from_numpy(np.vstack([e.state for e in experiences if e is not None]))
        actions = torch.from_numpy(np.vstack([e.action for e in experiences if e is not None]))
        rewards = torch.from_numpy(np.vstack([e.reward for e in experiences if e is not None]))
        next_states = torch.from_numpy(np.vstack([e.next_state for e in experiences if e is not None]))
        dones = torch.from_numpy(np.vstack([e.done for e in experiences if e is not None]))

        return (states, actions, rewards, next_states, dones)

    def __len__(self):
        """Return the current size of internal memory."""
        return len(self.memory)

```

```
In [20]: import random
```

```

import torch
import time

from collections import deque
import matplotlib.pyplot as plt
%matplotlib inline

##from ddpq_agent import Agent, ReplayBuffer

random_seed = random.randint(1,25)
#random_seed = 7
train_mode = True

agent = Agent(state_size=state_size, action_size=action_size, random_seed=random_seed)

In [21]: def ddpq(n_episodes=2000, max_t=1000, print_every=10, learn_every=20, num_learn=10, goal):
    total_scores_deque = deque(maxlen=100)
    total_scores = []

    for i_episode in range(1, n_episodes+1):
        #Reset the env and the agent
        env_info = env.reset(train_mode=train_mode)[brain_name] # reset the environment
        states = env_info.vector_observations # get the current states
        scores = np.zeros(num_agents) # initialize the score for each agent
        agent.reset()

        start_time = time.time()

        for t in range(max_t):
            actions = agent.act(states)

            env_info = env.step(actions)[brain_name] # send all actions to t
            next_states = env_info.vector_observations # get next state (for e
            rewards = env_info.rewards # get reward (for each

            dones = env_info.local_done # see if episode finish

            for state, action, reward, next_state, done in zip(states, actions, rewards, next_states, dones):
                agent.step(state, action, reward, next_state, done) # send actions to t

            scores += env_info.rewards # update the score (for
            states = next_states # roll over states to m

            if t%learn_every == 0:
                for _ in range(num_learn):
                    agent.start_learn()

            if np.any(dones): # exit loop if episode

```

```

        break

    mean_score = np.mean(scores)
    min_score = np.min(scores)
    max_score = np.max(scores)
    total_scores_deque.append(mean_score)
    total_scores.append(mean_score)
    total_average_score = np.mean(total_scores_deque)
    duration = time.time() - start_time

    print('\rEpisode {} \t Total Average Score: {:.2f} \t Mean: {:.2f} \t Min: {:.2f} \t Max: {:.2f}'.format(i_episode, total_average_score, mean_score, min_score, max_score, duration))

    if i_episode % print_every == 0:
        torch.save(agent.actor_local.state_dict(), 'checkpoint_actor.pth')
        torch.save(agent.critic_local.state_dict(), 'checkpoint_critic.pth')
        print('\rEpisode {} \t Total Average Score: {:.2f}'.format(i_episode, total_average_score))

    if total_average_score >= goal_score and i_episode >= 100:
        print('Problem Solved after {} epsisodes!! Total Average score: {:.2f}'.format(i_episode, total_average_score))
        torch.save(agent.actor_local.state_dict(), 'checkpoint_actor.pth')
        torch.save(agent.critic_local.state_dict(), 'checkpoint_critic.pth')
        break

    return total_scores

scores = ddpg()

```

Episode 1	Total Average Score: 0.80	Mean: 0.80	Min: 0.19	Max: 2.43
Episode 2	Total Average Score: 0.76	Mean: 0.72	Min: 0.01	Max: 1.64
Episode 3	Total Average Score: 0.78	Mean: 0.82	Min: 0.11	Max: 2.41
Episode 4	Total Average Score: 0.74	Mean: 0.60	Min: 0.00	Max: 1.56
Episode 5	Total Average Score: 0.77	Mean: 0.89	Min: 0.12	Max: 1.57
Episode 6	Total Average Score: 0.73	Mean: 0.56	Min: 0.11	Max: 1.56
Episode 7	Total Average Score: 0.77	Mean: 0.97	Min: 0.14	Max: 3.52
Episode 8	Total Average Score: 0.79	Mean: 0.99	Min: 0.00	Max: 2.94
Episode 9	Total Average Score: 0.85	Mean: 1.29	Min: 0.38	Max: 2.54
Episode 10	Total Average Score: 0.87	Mean: 1.09	Min: 0.09	Max: 1.92
Episode 10	Total Average Score: 0.87			
Episode 11	Total Average Score: 0.90	Mean: 1.20	Min: 0.11	Max: 2.13
Episode 12	Total Average Score: 0.91	Mean: 1.04	Min: 0.00	Max: 2.07
Episode 13	Total Average Score: 0.93	Mean: 1.09	Min: 0.00	Max: 2.96
Episode 14	Total Average Score: 0.91	Mean: 0.74	Min: 0.00	Max: 2.58
Episode 15	Total Average Score: 0.92	Mean: 1.02	Min: 0.06	Max: 3.40
Episode 16	Total Average Score: 0.95	Mean: 1.36	Min: 0.00	Max: 3.64
Episode 17	Total Average Score: 0.99	Mean: 1.59	Min: 0.34	Max: 2.77
Episode 18	Total Average Score: 1.03	Mean: 1.77	Min: 0.87	Max: 3.28
Episode 19	Total Average Score: 1.04	Mean: 1.30	Min: 0.31	Max: 2.51

Episode 20	Total Average Score: 1.09	Mean: 2.00	Min: 0.52	Max: 5.56
Episode 20	Total Average Score: 1.09			
Episode 21	Total Average Score: 1.11	Mean: 1.52	Min: 0.29	Max: 3.41
Episode 22	Total Average Score: 1.14	Mean: 1.71	Min: 0.53	Max: 2.97
Episode 23	Total Average Score: 1.17	Mean: 1.83	Min: 0.91	Max: 3.91
Episode 24	Total Average Score: 1.19	Mean: 1.69	Min: 0.90	Max: 3.35
Episode 25	Total Average Score: 1.23	Mean: 2.15	Min: 0.69	Max: 5.68
Episode 26	Total Average Score: 1.28	Mean: 2.55	Min: 0.83	Max: 6.22
Episode 27	Total Average Score: 1.32	Mean: 2.49	Min: 0.92	Max: 4.04
Episode 28	Total Average Score: 1.38	Mean: 2.84	Min: 0.86	Max: 5.23
Episode 29	Total Average Score: 1.39	Mean: 1.62	Min: 0.40	Max: 2.90
Episode 30	Total Average Score: 1.43	Mean: 2.72	Min: 0.16	Max: 6.47
Episode 30	Total Average Score: 1.43			
Episode 31	Total Average Score: 1.49	Mean: 3.15	Min: 1.02	Max: 6.92
Episode 32	Total Average Score: 1.54	Mean: 3.33	Min: 1.00	Max: 6.32
Episode 33	Total Average Score: 1.61	Mean: 3.68	Min: 1.06	Max: 6.39
Episode 34	Total Average Score: 1.67	Mean: 3.54	Min: 1.09	Max: 6.83
Episode 35	Total Average Score: 1.73	Mean: 3.93	Min: 0.65	Max: 6.49
Episode 36	Total Average Score: 1.80	Mean: 4.06	Min: 2.52	Max: 5.76
Episode 37	Total Average Score: 1.87	Mean: 4.36	Min: 1.38	Max: 7.91
Episode 38	Total Average Score: 1.93	Mean: 4.19	Min: 2.02	Max: 6.92
Episode 39	Total Average Score: 1.99	Mean: 4.30	Min: 1.72	Max: 8.69
Episode 40	Total Average Score: 2.05	Mean: 4.38	Min: 1.45	Max: 7.76
Episode 40	Total Average Score: 2.05			
Episode 41	Total Average Score: 2.12	Mean: 5.17	Min: 2.14	Max: 9.13
Episode 42	Total Average Score: 2.23	Mean: 6.68	Min: 3.30	Max: 10.58
Episode 43	Total Average Score: 2.29	Mean: 4.64	Min: 1.31	Max: 7.83
Episode 44	Total Average Score: 2.37	Mean: 6.01	Min: 2.76	Max: 11.34
Episode 45	Total Average Score: 2.46	Mean: 6.25	Min: 3.09	Max: 11.42
Episode 46	Total Average Score: 2.57	Mean: 7.78	Min: 4.22	Max: 12.64
Episode 47	Total Average Score: 2.66	Mean: 6.67	Min: 2.12	Max: 12.14
Episode 48	Total Average Score: 2.76	Mean: 7.19	Min: 3.96	Max: 10.08
Episode 49	Total Average Score: 2.85	Mean: 7.47	Min: 3.28	Max: 12.57
Episode 50	Total Average Score: 2.96	Mean: 8.18	Min: 5.10	Max: 12.96
Episode 50	Total Average Score: 2.96			
Episode 51	Total Average Score: 3.06	Mean: 8.24	Min: 6.10	Max: 12.28
Episode 52	Total Average Score: 3.15	Mean: 7.88	Min: 4.93	Max: 10.31
Episode 53	Total Average Score: 3.26	Mean: 8.63	Min: 4.72	Max: 14.63
Episode 54	Total Average Score: 3.37	Mean: 9.32	Min: 5.58	Max: 12.65
Episode 55	Total Average Score: 3.49	Mean: 10.06	Min: 3.23	Max: 14.4
Episode 56	Total Average Score: 3.58	Mean: 8.31	Min: 4.39	Max: 12.68
Episode 57	Total Average Score: 3.66	Mean: 8.31	Min: 3.95	Max: 12.41
Episode 58	Total Average Score: 3.76	Mean: 9.30	Min: 4.43	Max: 12.72
Episode 59	Total Average Score: 3.86	Mean: 9.70	Min: 5.45	Max: 14.95
Episode 60	Total Average Score: 3.97	Mean: 10.64	Min: 6.13	Max: 16.0
Episode 60	Total Average Score: 3.97			
Episode 61	Total Average Score: 4.09	Mean: 10.89	Min: 5.81	Max: 14.0
Episode 62	Total Average Score: 4.22	Mean: 12.64	Min: 5.78	Max: 18.6

Episode 63	Total Average Score: 4.33	Mean: 11.09	Min: 3.52	Max: 17.6
Episode 64	Total Average Score: 4.43	Mean: 10.51	Min: 3.79	Max: 15.2
Episode 65	Total Average Score: 4.53	Mean: 10.95	Min: 3.34	Max: 15.2
Episode 66	Total Average Score: 4.66	Mean: 13.34	Min: 7.82	Max: 17.4
Episode 67	Total Average Score: 4.78	Mean: 12.51	Min: 7.42	Max: 17.4
Episode 68	Total Average Score: 4.91	Mean: 13.92	Min: 9.58	Max: 19.1
Episode 69	Total Average Score: 5.04	Mean: 13.25	Min: 8.44	Max: 21.1
Episode 70	Total Average Score: 5.17	Mean: 14.62	Min: 8.41	Max: 18.7
Episode 70	Total Average Score: 5.17			
Episode 71	Total Average Score: 5.32	Mean: 15.44	Min: 10.01	Max: 23.
Episode 72	Total Average Score: 5.45	Mean: 15.19	Min: 9.29	Max: 21.9
Episode 73	Total Average Score: 5.56	Mean: 13.49	Min: 8.71	Max: 19.8
Episode 74	Total Average Score: 5.69	Mean: 15.01	Min: 9.54	Max: 23.9
Episode 75	Total Average Score: 5.82	Mean: 15.56	Min: 8.60	Max: 22.0
Episode 76	Total Average Score: 5.98	Mean: 17.85	Min: 9.57	Max: 27.9
Episode 77	Total Average Score: 6.15	Mean: 19.10	Min: 8.68	Max: 27.4
Episode 78	Total Average Score: 6.30	Mean: 18.04	Min: 7.98	Max: 26.8
Episode 79	Total Average Score: 6.41	Mean: 14.35	Min: 7.41	Max: 22.7
Episode 80	Total Average Score: 6.53	Mean: 16.22	Min: 5.08	Max: 27.8
Episode 80	Total Average Score: 6.53			
Episode 81	Total Average Score: 6.68	Mean: 19.20	Min: 11.40	Max: 28.
Episode 82	Total Average Score: 6.83	Mean: 18.83	Min: 12.85	Max: 29.
Episode 83	Total Average Score: 6.99	Mean: 19.92	Min: 11.80	Max: 27.
Episode 84	Total Average Score: 7.17	Mean: 21.82	Min: 8.97	Max: 39.1
Episode 85	Total Average Score: 7.33	Mean: 20.68	Min: 12.41	Max: 29.
Episode 86	Total Average Score: 7.50	Mean: 21.96	Min: 11.37	Max: 31.
Episode 87	Total Average Score: 7.66	Mean: 21.67	Min: 12.82	Max: 27.
Episode 88	Total Average Score: 7.83	Mean: 22.82	Min: 14.61	Max: 29.
Episode 89	Total Average Score: 7.98	Mean: 21.46	Min: 16.03	Max: 29.
Episode 90	Total Average Score: 8.15	Mean: 22.86	Min: 15.89	Max: 31.
Episode 90	Total Average Score: 8.15			
Episode 91	Total Average Score: 8.32	Mean: 23.99	Min: 16.12	Max: 38.
Episode 92	Total Average Score: 8.51	Mean: 25.71	Min: 13.08	Max: 39.
Episode 93	Total Average Score: 8.69	Mean: 24.60	Min: 16.72	Max: 35.
Episode 94	Total Average Score: 8.87	Mean: 26.38	Min: 18.28	Max: 37.
Episode 95	Total Average Score: 9.05	Mean: 25.36	Min: 14.94	Max: 31.
Episode 96	Total Average Score: 9.22	Mean: 26.02	Min: 19.93	Max: 30.
Episode 97	Total Average Score: 9.41	Mean: 27.11	Min: 21.53	Max: 31.
Episode 98	Total Average Score: 9.58	Mean: 26.55	Min: 21.78	Max: 33.
Episode 99	Total Average Score: 9.76	Mean: 27.44	Min: 21.11	Max: 34.
Episode 100	Total Average Score: 9.95	Mean: 28.19	Min: 19.54	Max: 37
Episode 100	Total Average Score: 9.95			
Episode 101	Total Average Score: 10.22	Mean: 28.16	Min: 19.90	Max: 3
Episode 102	Total Average Score: 10.50	Mean: 28.07	Min: 23.21	Max: 3
Episode 103	Total Average Score: 10.77	Mean: 27.99	Min: 23.93	Max: 3
Episode 104	Total Average Score: 11.05	Mean: 28.63	Min: 23.26	Max: 3
Episode 105	Total Average Score: 11.33	Mean: 28.82	Min: 11.94	Max: 3
Episode 106	Total Average Score: 11.63	Mean: 30.49	Min: 24.58	Max: 3


Episode 107	Total Average Score: 11.91	Mean: 29.20	Min: 20.82	Max: 3
Episode 108	Total Average Score: 12.19	Mean: 28.89	Min: 23.98	Max: 3
Episode 109	Total Average Score: 12.48	Mean: 30.98	Min: 27.25	Max: 3
Episode 110	Total Average Score: 12.77	Mean: 29.68	Min: 24.00	Max: 3
Episode 110	Total Average Score: 12.77			
Episode 111	Total Average Score: 13.06	Mean: 30.52	Min: 24.51	Max: 3
Episode 112	Total Average Score: 13.37	Mean: 31.23	Min: 24.61	Max: 3
Episode 113	Total Average Score: 13.69	Mean: 33.05	Min: 29.27	Max: 3
Episode 114	Total Average Score: 13.99	Mean: 31.65	Min: 22.70	Max: 3
Episode 115	Total Average Score: 14.30	Mean: 31.31	Min: 14.36	Max: 3
Episode 116	Total Average Score: 14.61	Mean: 32.70	Min: 28.65	Max: 3
Episode 117	Total Average Score: 14.92	Mean: 32.23	Min: 26.77	Max: 3
Episode 118	Total Average Score: 15.23	Mean: 33.05	Min: 28.82	Max: 3
Episode 119	Total Average Score: 15.54	Mean: 32.46	Min: 24.74	Max: 3
Episode 120	Total Average Score: 15.84	Mean: 31.60	Min: 22.82	Max: 3
Episode 120	Total Average Score: 15.84			
Episode 121	Total Average Score: 16.14	Mean: 31.67	Min: 22.26	Max: 3
Episode 122	Total Average Score: 16.46	Mean: 34.15	Min: 28.22	Max: 3
Episode 123	Total Average Score: 16.78	Mean: 33.02	Min: 26.87	Max: 3
Episode 124	Total Average Score: 17.10	Mean: 34.62	Min: 30.68	Max: 3
Episode 125	Total Average Score: 17.42	Mean: 33.63	Min: 29.76	Max: 3
Episode 126	Total Average Score: 17.74	Mean: 34.22	Min: 28.36	Max: 3
Episode 127	Total Average Score: 18.07	Mean: 35.50	Min: 31.69	Max: 3
Episode 128	Total Average Score: 18.38	Mean: 34.26	Min: 31.54	Max: 3
Episode 129	Total Average Score: 18.72	Mean: 35.38	Min: 32.86	Max: 3
Episode 130	Total Average Score: 19.04	Mean: 34.55	Min: 26.34	Max: 3
Episode 130	Total Average Score: 19.04			
Episode 131	Total Average Score: 19.36	Mean: 35.44	Min: 31.65	Max: 3
Episode 132	Total Average Score: 19.67	Mean: 34.58	Min: 30.84	Max: 3
Episode 133	Total Average Score: 19.99	Mean: 35.42	Min: 32.91	Max: 3
Episode 134	Total Average Score: 20.31	Mean: 35.48	Min: 30.98	Max: 3
Episode 135	Total Average Score: 20.63	Mean: 36.30	Min: 30.54	Max: 3
Episode 136	Total Average Score: 20.94	Mean: 35.22	Min: 31.97	Max: 3
Episode 137	Total Average Score: 21.26	Mean: 35.70	Min: 33.55	Max: 3
Episode 138	Total Average Score: 21.58	Mean: 36.39	Min: 31.93	Max: 3
Episode 139	Total Average Score: 21.91	Mean: 36.90	Min: 33.94	Max: 3
Episode 140	Total Average Score: 22.23	Mean: 36.54	Min: 32.31	Max: 3
Episode 140	Total Average Score: 22.23			
Episode 141	Total Average Score: 22.53	Mean: 35.50	Min: 29.89	Max: 3
Episode 142	Total Average Score: 22.81	Mean: 35.07	Min: 31.72	Max: 3
Episode 143	Total Average Score: 23.13	Mean: 36.52	Min: 32.19	Max: 3
Episode 144	Total Average Score: 23.43	Mean: 35.63	Min: 31.69	Max: 3
Episode 145	Total Average Score: 23.72	Mean: 35.60	Min: 31.95	Max: 3
Episode 146	Total Average Score: 24.01	Mean: 36.16	Min: 33.58	Max: 3
Episode 147	Total Average Score: 24.31	Mean: 37.30	Min: 35.45	Max: 3
Episode 148	Total Average Score: 24.59	Mean: 34.73	Min: 22.48	Max: 3
Episode 149	Total Average Score: 24.87	Mean: 35.84	Min: 32.22	Max: 3
Episode 150	Total Average Score: 25.17	Mean: 37.62	Min: 34.45	Max: 3

Episode 150	Total Average Score: 25.17			
Episode 151	Total Average Score: 25.46	Mean: 37.41	Min: 34.91	Max: 3
Episode 152	Total Average Score: 25.75	Mean: 37.57	Min: 33.36	Max: 3
Episode 153	Total Average Score: 26.04	Mean: 37.28	Min: 32.34	Max: 3
Episode 154	Total Average Score: 26.30	Mean: 35.56	Min: 32.22	Max: 3
Episode 155	Total Average Score: 26.57	Mean: 37.05	Min: 35.31	Max: 3
Episode 156	Total Average Score: 26.85	Mean: 35.85	Min: 34.31	Max: 3
Episode 157	Total Average Score: 27.13	Mean: 36.39	Min: 33.80	Max: 3
Episode 158	Total Average Score: 27.42	Mean: 37.85	Min: 35.58	Max: 3
Episode 159	Total Average Score: 27.69	Mean: 36.95	Min: 33.52	Max: 3
Episode 160	Total Average Score: 27.96	Mean: 37.79	Min: 34.55	Max: 3
Episode 160	Total Average Score: 27.96			
Episode 161	Total Average Score: 28.23	Mean: 38.46	Min: 36.21	Max: 3
Episode 162	Total Average Score: 28.49	Mean: 37.75	Min: 35.65	Max: 3
Episode 163	Total Average Score: 28.75	Mean: 37.95	Min: 36.79	Max: 3
Episode 164	Total Average Score: 29.02	Mean: 37.28	Min: 34.44	Max: 3
Episode 165	Total Average Score: 29.29	Mean: 37.75	Min: 32.75	Max: 3
Episode 166	Total Average Score: 29.50	Mean: 34.45	Min: 28.86	Max: 3
Episode 167	Total Average Score: 29.75	Mean: 37.14	Min: 33.97	Max: 3
Episode 168	Total Average Score: 29.98	Mean: 37.56	Min: 36.11	Max: 3
Episode 169	Total Average Score: 30.22	Mean: 36.90	Min: 34.74	Max: 3

Problem Solved after 169 epsisodes!! Total Average score: 30.22

```
In [22]: fig = plt.figure()
         ax = fig.add_subplot(1,1,1)

         plt.plot(np.arange(1, len(scores)+1), scores)
         plt.ylabel('Score')
         plt.xlabel('Episode #')
         plt.show()
```



```
In [23]: # Load the saved weights into Pytorch model
agent.actor_local.load_state_dict(torch.load('checkpoint_actor.pth', map_location='cpu'))
agent.critic_local.load_state_dict(torch.load('checkpoint_critic.pth', map_location='cpu'))

env_info = env.reset(train_mode=False)[brain_name]           # reset the environment
states = env_info.vector_observations                        # get the current state (for each agent)
scores = np.zeros(num_agents)                               # initialize the score (for each agent)

while True:
    actions = agent.act(states)                               # select actions from loaded model
    env_info = env.step(actions)[brain_name]                 # send all actions to the environment
    next_states = env_info.vector_observations                # get next state (for each agent)
    rewards = env_info.rewards                               # get reward (for each agent)
    dones = env_info.local_done                             # see if episode finished
    scores += env_info.rewards                               # update the score (for each agent)
    states = next_states                                     # roll over states to next time
    if np.any(dones):                                       # exit loop if episode finished
        break
    print('Total score: {}'.format(np.mean(scores)))

Total score: 36.58749918220565

In [24]: env.close()
```